

Introduction to supervised machine learning, k-fold cross-validation, nearest neighbors, and linear models

Toby Dylan Hocking

Supervised machine learning

- ▶ Goal is to learn a function $f(\mathbf{x}) = y$ where \mathbf{x} is an input/feature vector and y is an output/label.
- ▶ x = image of digit/clothing, $y \in \{0, \dots, 9\}$ (ten classes).
- ▶ x = vector of word counts in email, $y \in \{1, 0\}$ (spam or not).
- ▶ x = image of retina, y = risk score for heart disease.
- ▶ This week we will focus on a specific kind of supervised learning problem called binary classification, which means $y \in \{1, 0\}$.

Learning algorithm

- ▶ We want a learning algorithm `LEARN` which inputs a training data set and outputs a prediction function f .
- ▶ In math a training data set with n observations and p features is a matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ with a label vector $\mathbf{y} \in \{0, 1\}^n$.
- ▶ On computers it is a CSV file with n rows and $p + 1$ columns.
- ▶ Want: $\text{LEARN}(\mathbf{X}, \mathbf{y}) \rightarrow f$.
- ▶ We will use three such data sets from Elements of Statistical Learning book by Hastie et al. (mixture slightly modified)

name	observations, n	inputs/features, p	outputs/labels
zip.test	images, 623	pixel intensities, 256	0/1 digits
spam	emails, 4601	word counts, 57	spam=1/not=0
mixture	people, 200	height/weight, 2	democratic/republican

<https://github.com/tdhock/cs570-spring-2022/tree/master/data>

<https://hastie.su.domains/ElemStatLearn/data.html>

Mixture data table

```
##           party  height_in  weight_lb
## 0    democratic  71.741421  149.565034
## 1    democratic  69.582283  149.275446
## 2    democratic  69.983547  149.961470
## 3    democratic  69.908764  150.021178
## 4    democratic  69.195491  150.111237
## ..          ...          ...          ...
## 195 republican  69.472078  151.537588
## 196 republican  71.140501  149.409036
## 197 republican  70.517269  150.236183
## 198 republican  69.223459  151.486248
## 199 republican  69.019082  149.795387
##
## [200 rows x 3 columns]
```

Spam data table

##	0	1	2	...	55	56	57
## 0	0.00	0.64	0.64	...	61	278	1
## 1	0.21	0.28	0.50	...	101	1028	1
## 2	0.06	0.00	0.71	...	485	2259	1
## 3	0.00	0.00	0.00	...	40	191	1
## 4	0.00	0.00	0.00	...	40	191	1
##
## 4596	0.31	0.00	0.62	...	3	88	0
## 4597	0.00	0.00	0.00	...	4	14	0
## 4598	0.30	0.00	0.30	...	6	118	0
## 4599	0.96	0.00	0.00	...	5	78	0
## 4600	0.00	0.00	0.65	...	5	40	0

##

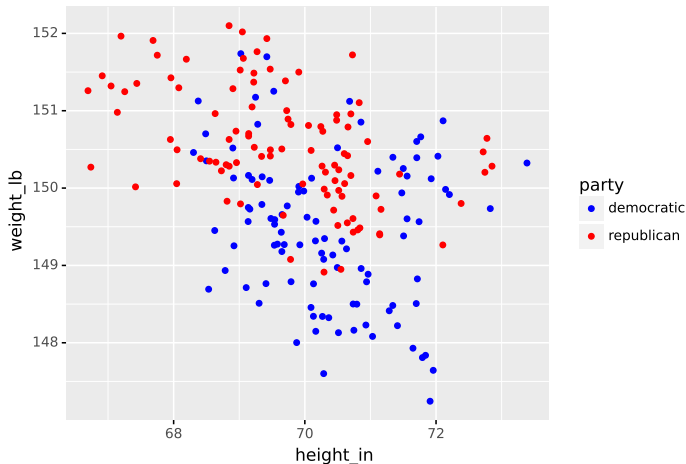
[4601 rows x 58 columns]

Zip.test data table

```
##           0      1      2      ...  254  255  256
## 0          9 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 1          6 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 2          3 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 3          6 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 4          6 -1.0 -1.0      ... -1.0 -1.0 -1.0
## ...      ...   ...   ...   ...   ...   ...   ...
## 2002       3 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 2003       9 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 2004       4 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 2005       0 -1.0 -1.0      ... -1.0 -1.0 -1.0
## 2006       1 -1.0 -1.0      ... -1.0 -1.0 -1.0
##
## [2007 rows x 257 columns]
```

Visualize mixture data set


- ▶ Each axis represents one column of the \mathbf{X} matrix.
- ▶ Each point represents one row of the \mathbf{X} matrix.
- ▶ Color represents class label \mathbf{y} .




A basic machine learning algorithm

- ▶ Goal of supervised learning is to learn a function which predicts the label for new inputs $x \in \mathbb{R}^2$.
- ▶ K-Nearest neighbors: a simple non-linear algorithm.
- ▶ For any new data point, predict the average label of the K nearest neighbors.

Visualize predictions of 1-nearest neighbor algorithm



Also plot decision boundary in black

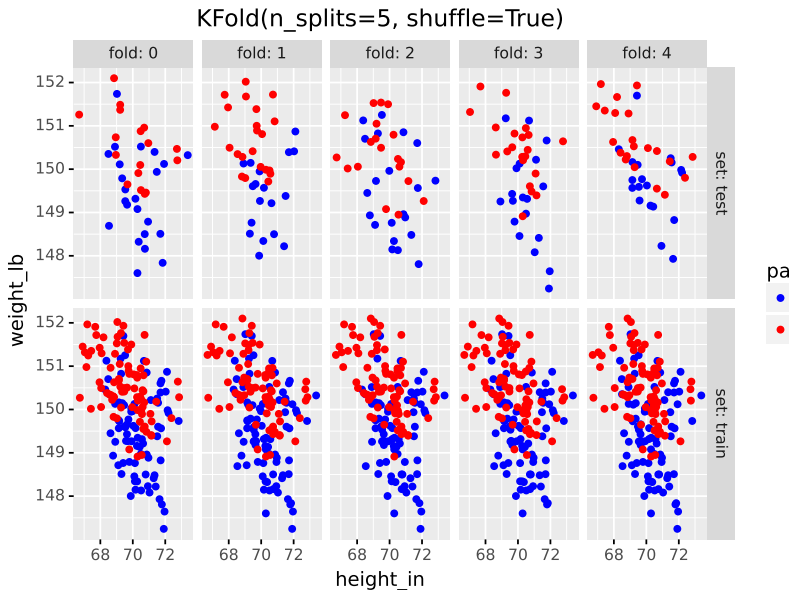


Is it good to have 100% accuracy on train data?

- ▶ Remember: goal is function f with accurate predictions on new inputs.
- ▶ What is a new input?
- ▶ We must assume that new/test inputs are similar to old/train inputs.
- ▶ In the statistical literature this is the iid (independent and identically distributed) assumption.
- ▶ We can therefore split the full data set into train/test sets.
- ▶ Train set is used to learn the prediction function f .
- ▶ Test set (simulated new inputs) is used to evaluate the accuracy of the function f (but can not be used to learn function f).

K-fold cross-validation for splitting data

Visualization of splits/sets in input/feature space



Basic idea of nearest neighbors

Basic idea of linear model

Visualize iris data without labels

- ▶ Let $X \in \mathbb{R}^{150 \times 2}$ be the data matrix (input for clustering).