

Introduction to supervised machine learning,
k-fold cross-validation, nearest neighbors, and
linear models

Toby Dylan Hocking

Supervised machine learning

- ▶ Goal is to learn a function $f(\mathbf{x}) = y$ where \mathbf{x} is an input/feature vector and y is an output/label.
- ▶ x = image of digit/clothing, $y \in \{0, \dots, 9\}$ (ten classes).
- ▶ x = vector of word counts in email, $y \in \{1, 0\}$ (spam or not).
- ▶ x = image of retina, y = risk score for heart disease.
- ▶ This week we will focus on a specific kind of supervised learning problem called binary classification, which means $y \in \{1, 0\}$.

Learning algorithm

- ▶ We want a learning algorithm `LEARN` which inputs a training data set and outputs a prediction function f .
- ▶ In math a training data set with n observations and p features is a matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ with a label vector $\mathbf{y} \in \{0, 1\}^n$.
- ▶ On computers it is a CSV file with n rows and $p + 1$ columns.
- ▶ Want: $\text{LEARN}(\mathbf{X}, \mathbf{y}) \rightarrow f$.
- ▶ We will use three such data sets from Elements of Statistical Learning book by Hastie et al. (mixture slightly modified)

name	observations, n	inputs/features, p	outputs/labels
zip.test	images, 623	pixel intensities, 256	0/1 digits
spam	emails, 4601	word counts, 57	spam=1/not=0
mixture	people, 200	height/weight, 2	democratic/republican

<https://github.com/tdhock/cs570-spring-2022/tree/master/data>

<https://hastie.su.domains/ElemStatLearn/data.html>

Mixture data table

```
##           party height_in  weight_lb
## 0    democratic  71.741421  149.565034
## 1    democratic  69.582283  149.275446
## 2    democratic  69.983547  149.961470
## 3    democratic  69.908764  150.021178
## 4    democratic  69.195491  150.111237
## ..          ...          ...          ...
## 195 republican  69.472078  151.537588
## 196 republican  71.140501  149.409036
## 197 republican  70.517269  150.236183
## 198 republican  69.223459  151.486248
## 199 republican  69.019082  149.795387
##
## [200 rows x 3 columns]
```

Spam data table


```
##           0         1         2     ...     55      56    57
## 0      0.00    0.64    0.64    ...     61     278     1
## 1      0.21    0.28    0.50    ...    101    1028     1
## 2      0.06    0.00    0.71    ...    485    2259     1
## 3      0.00    0.00    0.00    ...     40     191     1
## 4      0.00    0.00    0.00    ...     40     191     1
## ...      ...      ...      ...     ...     ...     ..
## 4596    0.31    0.00    0.62    ...      3      88      0
## 4597    0.00    0.00    0.00    ...      4      14      0
## 4598    0.30    0.00    0.30    ...      6     118      0
## 4599    0.96    0.00    0.00    ...      5      78      0
## 4600    0.00    0.00    0.65    ...      5      40      0
##
## [4601 rows x 58 columns]
```

Zip.test data table

```
##          0      1      2      ...    254    255    256
## 0         9 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 1         6 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 2         3 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 3         6 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 4         6 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## ...      ...    ...    ...    ...    ...    ...    ...
## 2002       3 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 2003       9 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 2004       4 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 2005       0 -1.0 -1.0      ...   -1.0  -1.0  -1.0
## 2006       1 -1.0 -1.0      ...   -1.0  -1.0  -1.0
##
## [2007 rows x 257 columns]
```

Visualize mixture data set

- ▶ Each axis represents one column of the \mathbf{X} matrix.
- ▶ Each point represents one row of the \mathbf{X} matrix.
- ▶ Color represents class label \mathbf{y} .



function viz

Visualize iris data without labels

- ▶ Let $X \in \mathbb{R}^{150 \times 2}$ be the data matrix (input for clustering).